

Assignment 3: Q-Learning

In this assignment you will implement deep Q-networks (DQN) [1], a Q-learning algorithm that leverages deep neural networks, to play Atari games. The template code is available at:

https://github.com/xbpeng/rl_assignments

Installation instructions are provided in `README.md`. All of the files that need to be changed in this assignment are located in the `a3/` directory. No files outside of this directory should be modified. Locations where code needs to be modified are labeled `TODO`.

1 Q-Learning

Q-learning is a family of reinforcement learning algorithms that solves an MDP by learning a Q-function $Q(\mathbf{s}, \mathbf{a})$, instead of directly learning a policy $\pi(\mathbf{a}|\mathbf{s})$. The Q-function Q^π provides an estimate of the expected return of performing an action \mathbf{a} in state \mathbf{s} , and following a policy π for all future timesteps,

$$Q^\pi(\mathbf{s}, \mathbf{a}) = \mathbb{E}_{\tau \sim p(\tau|\pi, \mathbf{s}_0=\mathbf{s}, \mathbf{a}_0=\mathbf{a})} \left[\sum_t \gamma^t r_t \right]. \quad (1)$$

The optimal Q-function for a given MDP can be learned using an iterative fixed-point method, where at each iteration k , a new Q-function Q^{k+1} is constructed via the Bellman update by bootstrapping from the current Q-function Q^k :

$$Q^{k+1}(\mathbf{s}, \mathbf{a}) = \mathbb{E}_{\mathbf{s}' \sim p(\mathbf{s}'|\mathbf{s}, \mathbf{a})} \left[r(\mathbf{s}, \mathbf{a}, \mathbf{s}') + \gamma \max_{\mathbf{a}'} Q^k(\mathbf{s}', \mathbf{a}') \right], \quad (2)$$

where $r(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ is the reward function, $\gamma \in [0, 1]$ is a discount factor, and $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$ represents the dynamics of the environment.

A policy can be recovered from a given Q-function by selecting the action with the maximum predicted Q-value at a given state:

$$\pi^k(\mathbf{a}|\mathbf{s}) = \begin{cases} 1 & \text{if } \mathbf{a} = \arg \max_{\mathbf{a}'} Q^k(\mathbf{s}, \mathbf{a}') \\ 0 & \text{otherwise} \end{cases}. \quad (3)$$

This arg-max procedure returns a deterministic policy that selects actions greedily according to the Q-function. If Q-function is the optimal Q-function Q^* , then the resulting arg-max policy will be an optimal policy π^* . However, if the Q-function is not optimal, then a deterministic greedy policy will often lead to insufficient exploration during the learning process. ϵ -greedy exploration is a simple strategy that mitigate this exploration issue by constructing a stochastic policy

$$\pi^k(\mathbf{a}|\mathbf{s}) = \begin{cases} 1 - \epsilon & \text{if } \mathbf{a} = \arg \max_{\mathbf{a}'} Q^k(\mathbf{s}, \mathbf{a}') \\ \epsilon & \text{otherwise} \end{cases}, \quad (4)$$

which greedily select the best action according to the current Q-function with probability $1 - \epsilon$, and selects a random action with probability ϵ . This method helps to ensure that the agent can explore new actions and potentially discover more optimal strategies during the training process. Pseudo-code for Q-learning is available in Algorithm 1.

ALGORITHM 1: Q-Learning

- 1: $Q^0 \leftarrow$ initialize Q-function
 - 2: $\mathcal{D} \leftarrow \{\emptyset\}$ initialize dataset
 - 3: **for** iteration $k = 0, \dots, n - 1$ **do**
 - 4: Sample trajectory τ according to $Q^k(\mathbf{s}, \mathbf{a})$
 - 5: Add transitions to dataset $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
 - 6: Calculate target values for each sample i :
 $y_i = r_i + \gamma \max_{\mathbf{a}'} Q^k(\mathbf{s}'_i, \mathbf{a}')$
 - 7: Update Q-function:
 $Q^{k+1} = \arg \min_Q \mathbb{E}_{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i) \sim \mathcal{D}} [(y_i - Q(\mathbf{s}_i, \mathbf{a}_i))^2]$
 - 8: **end for**
 - 9: return Q^n
-

In this assignment, you will be implementing the deep Q-networks (DQN) algorithm to play Atari games [1]. To run DQN on the **Pong** task, use the following command:

```
python run.py --mode train \
--env_config data/envs/atari_pong.yaml \
--agent_config a3/atari_dqn_agent.yaml \
--log_file output/log.txt \
--out_model_file output/model.pt \
--max_samples 3000000 \
--visualize
```

`--visualize` should be disabled for faster training. Once a model has been trained, you can load a checkpoint and test the model with the following command:

```
python run.py --mode test \
--env_config data/envs/atari_pong.yaml \
--agent_config a3/atari_dqn_agent.yaml \
--model_file output/model.pt \
--test_episodes 20 \
--visualize
```

1.1 Epsilon-Greedy Exploration

In `a3/dqn_agent.py`, implement the `_get_exp_prob()` method, which calculates the probability ϵ of selecting a random action in ϵ -greedy exploration. Since the Q-function is randomly initialized at the start of training, the exploration probability typically starts with a high value (e.g. $\epsilon = 1$), and then decreases to a low value (e.g. $\epsilon = 0.1$) over the course of training as the Q-function improves. The exploration probability starts with a value of ϵ_{beg} , and then linearly annealed to ϵ_{end} over the course of n_{max} samples,

$$\epsilon(n) = (1 - l)\epsilon_{\text{beg}} + l\epsilon_{\text{end}}, \quad l = \text{clip}\left(\frac{n}{n_{\text{max}}}, 0, 1\right), \quad (5)$$

where n denotes the total number of samples collected for training so far. The initial exploration probability ϵ_{beg} is given by `self._exp_prob_beg`, and the final probability ϵ_{end} is given by `self._exp_prob_end`. The total number of samples n_{max} for annealing the exploration probability is specified by `self._exp_anneal_samples`, and the number of samples n collected so far is recorded in `self._sample_count`. The output should be a scalar value corresponding to the probability of selecting a random action.

1.2 Action Sampling

In `a3/dqn_agent.py`, implement the `_sample_action(qs)` method, which samples actions according to the Q-values of each action. The input consists of a tensor of `qs`, which contains the predicted Q-values of each action. Implement ϵ -greedy exploration (Equation 4), where the probability of sampling a random action ϵ is specified by `self._get_exp_prob()`. With probability $1 - \epsilon$, greedily select the action with the highest Q-value. With probability ϵ , select a random action uniformly from the set of possible actions. The output `a` should be a tensor containing the index of the selected action.

1.3 Target Values

In `a3/dqn_agent.py`, implement the `_compute_tar_vals(r, norm_next_obs, done)` method, which calculates target values for updating the Q-function. The inputs consist of a tensor of rewards `r`, normalized observations at the next timestep `norm_next_obs`, and done flags `done` indicating if a timestep is the last timestep of an episode. The target value y_i is then calculated according to

$$y_i = r_i + \gamma(1 - \text{done}_i) \max_{\mathbf{a}'_i} Q^{\text{tar}}(\mathbf{s}'_i, \mathbf{a}'_i). \quad (6)$$

If a sample corresponds to the last timestep of an episode (i.e. $\text{done}_i = 1$), then the Q-value of the next timestep should be set to 0.

One of the innovations of DQN is the use of a target model Q^{tar} to calculate the target values [1], instead of directly calculating target values using the current Q-function. The target model is a delayed version of the Q-function, which is a copy of the parameters of the main Q-function from a number of iterations ago. Q^{tar} is kept fixed for a number of iteration,

before being updated with a copy of the parameters from the latest Q-function. By keeping Q^{tar} fixed for a number of iterations, the target model provides more stable target values for updating the main Q-function. The target model is given by `self._tar_model`, and the main model is given by `self._model`. The Q-function can be queried by using the method `eval_q(norm_obs)`. The output `tar_vals` should be a tensor containing the target values for each sample.

1.4 Loss Calculation

In `a3/dqn_agent.py`, implement the `_compute_q_loss(norm_obs, a, tar_vals)` method, which compute the loss for updating the Q-function. The input consists of a tensor of normalized observations `norm_obs`, a tensor containing the indices of discrete actions selected at each timestep `a`, and target values for each timestep `tar_vals`. The output `loss` should be a scalar tensor containing the loss for updating the Q-function. The loss should be calculated as the mean squared-error between the target values and the outputs of the Q-function

$$l(Q) = \mathbb{E}_{\mathbf{s}_i, \mathbf{a}_i, y_i \sim \mathcal{D}} [(y_i - Q(\mathbf{s}_i, \mathbf{a}_i))^2]. \quad (7)$$

Note, the loss should only be applied to the main model `self._model`, and should not be applied to the target model `self._tar_model`.

1.5 Target Model Update

In `a3/dqn_agent.py`, implement the `_sync_tar_model()` method, which updates the target model by copying the parameters from the main model. The main model is given by `self._model`, and the target model is given by `self._tar_model`. This method is used to periodically synchronize the parameters of the target model and the main model. `self._model.parameters()` can be used to retrieve a list of tensors containing the parameters of a model.

1.6 Tasks

Train DQN policies to play two Atari games, `Pong` and `Breakout`. The two games can be specified using the environment config files `data/envs/atari_pong.yaml` and `data/envs/atari_breakout.yaml` respectively. Train each model for at least 3 million timesteps. Tune the hyperparameters in the agent config file `a3/atari_dqn_agent.yaml` so that the policies reach a test return of at least 19 for `Pong` and 40 for `Breakout`. Plot a learning curve of the test return of each model using the plotting script `tools/plot_log/plot_log.py`.

2 Bonus

1 bonus point will be award to the submission that achieves the best performance on the `Breakout` tasks using DQN. To improve the performance of DQN, you can do more ex-

tensive tuning of the hyperparameters, as well as make any changes to the algorithm in `dqn_agent.py`. Submit any modifications made to the DQN algorithm in a separate file `bonus_dqn_agent.py`, a checkpoint of the trained model `bonus_breakout_model.pt`, and a text file `bonus.txt` detailing the modifications you made.

Submission

Your submission should contain the following files:

- `dqn_agent.py`: code changes.
- `atari_dqn_agent.yaml`: tuned hyperparameters.
- `atari_pong_dqn_model.pt`: trained model for the `Pong` task.
- `atari_breakout_dqn_model.pt`: trained model for the `Breakout` task.
- `atari_pong_dqn_log.txt`: training log for the `Pong` task.
- `atari_breakout_dqn_log.txt`: training log for the `Breakout` task.
- `atari_pong_dqn_curve.png`: image of the learning curve for the `Pong` task.
- `atari_breakout_dqn_curve.png`: image of the learning curve for the `Breakout` task.
- `bonus_dqn_agent.py`: modified version of the DQN agent for the bonus component.
- `bonus_breakout_model.pt`: trained `Breakout` model for the bonus component.
- `bonus.txt`: a text file detailing any modifications you made in `bonus_dqn_agent.py` for improving the performance of the algorithm.

All files should be stored in a directory named `a3`, and then zip the directory for submission. Do not add any additional subdirectories.

References

- [1] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik, I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, Feb. 2015.